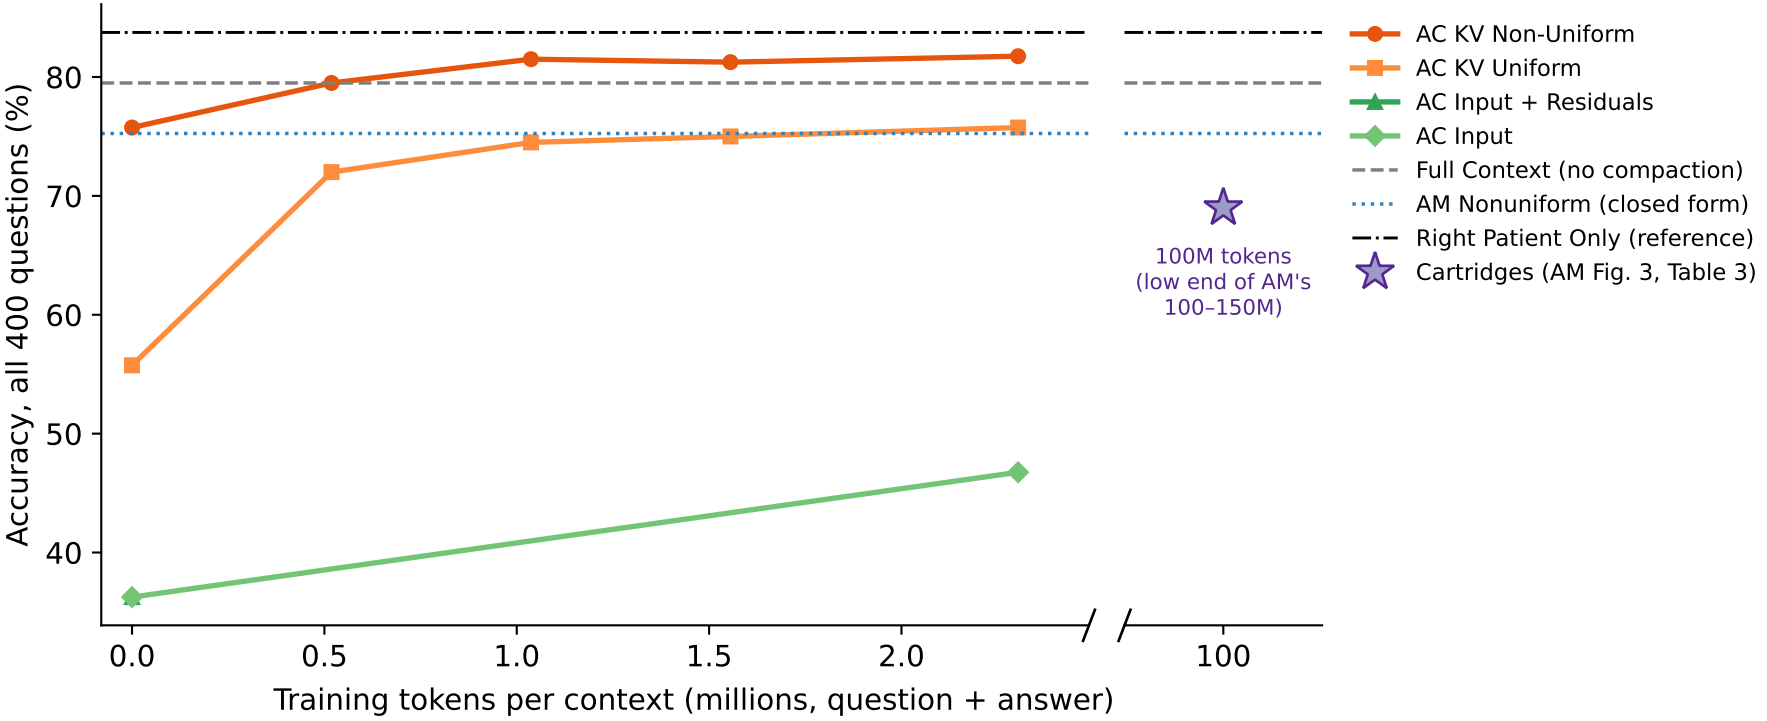
Accuracy, all 400 questions (%)
80
70
60
50
40
0.0
0.5
1.0
1.5
2.0
100
Training tokens per context (millions, question + answer)
100M tokens
(low end of AM's
100–150M)
AC KV Non-Uniform
AC KV Uniform
AC Input + Residuals
AC Input
Full Context (no compaction)
AM Nonuniform (closed form)
Right Patient Only (reference)
Cartridges (AM Fig. 3, Table 3)