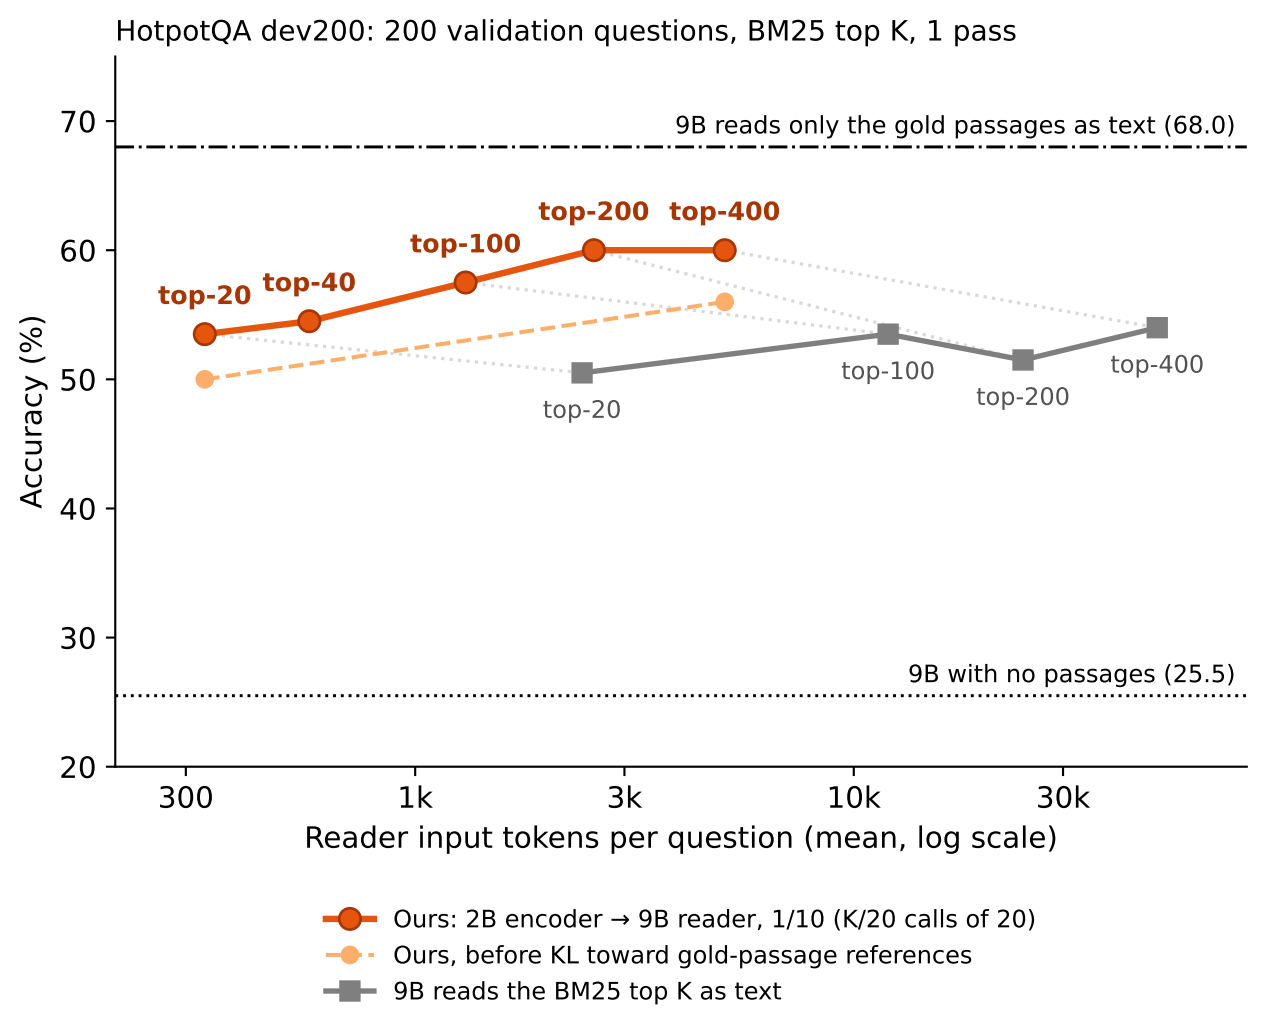
HotpotQA dev200: 200 validation questions, BM25 top K, 1 pass
9B reads only the gold passages as text (68.0)
9B with no passages (25.5)
top-20
top-40
top-100
top-200
top-400
top-20
top-100
top-200
top-400
Accuracy (%)
70
60
50
40
30
20
300
1k
3k
10k
30k
Reader input tokens per question (mean, log scale)
Ours: 2B encoder → 9B reader, 1/10 (K/20 calls of 20)
Ours, before KL toward gold-passage references
9B reads the BM25 top K as text