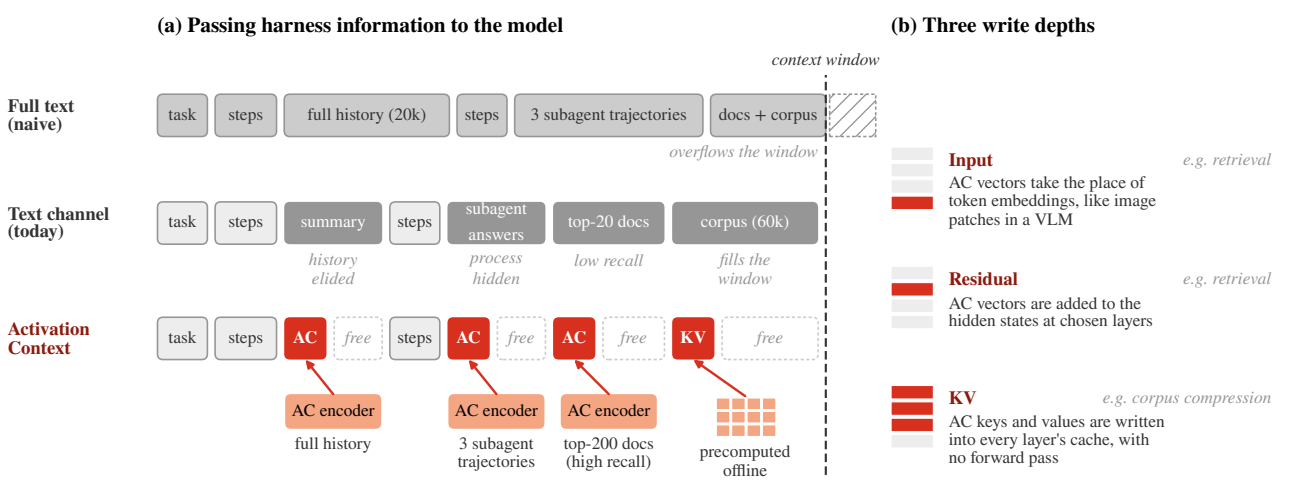
(a) Passing harness information to the model
(b) Three write depths
context window
Full text (naive)
task
steps
full history (20k)
steps
3 subagent trajectories
docs + corpus
overflows the window
Text channel (today)
task
steps
summary
steps
subagent answers
top-20 docs
corpus (60k)
history elided
process hidden
low recall
fills the window
Activation Context
task
steps
AC
free
steps
AC
free
AC
free
KV
free
AC encoder
full history
AC encoder
3 subagent trajectories
AC encoder
top-200 docs (high recall)
precomputed offline
Input
e.g. retrieval
AC vectors take the place of token embeddings, like image patches in a VLM
Residual
e.g. retrieval
AC vectors are added to the hidden states at chosen layers
KV
e.g. corpus compression
AC keys and values are written into every layer's cache, with no forward pass